

SQL:

Data Manipulation Language

csc343, Introduction to Databases

Diane Horton

Winter 2017



UNIVERSITY OF
TORONTO

Introduction

- So far, we have defined database schemas and queries mathematically.
- SQL is a formal language for doing so with a DBMS.
- “Structured Query Language”, but it’s for more than writing queries.
- Two sub-parts:
 - DDL (Data Definition Language), for defining schemas.
 - DML (Data Manipulation Language), for writing queries and modifying the database.

PostgreSQL

- We'll be working in PostgreSQL, an open-source relational DBMS.
- Learn your way around the documentation; it will be very helpful.
- Standards?
 - There are several, the most recent being SQL:2008.
 - The standards are not freely available. Must purchase from the International Standards Organization (ISO).
 - PostgreSQL supports most of it SQL:2008.
 - DBMSs vary in the details around the edges, making portability difficult.

A high-level language

- SQL is a very high-level language.
 - Say “what” rather than “how.”
- You write queries without manipulating data. Contrast languages like Java or C++.
- Provides physical “data independence:”
 - Details of how the data is stored can change with no impact on your queries.
- You can focus on readability.
 - But because the DMBS optimizes your query, you get efficiency.

Heads up: SELECT vs σ

- In SQL,
 - “SELECT” is for choosing columns, *i.e.*, Π .
 - Example:

```
SELECT surName
FROM Student
WHERE campus = 'StG';
```
- In relational algebra,
 - “select” means choosing rows, *i.e.*, σ .

Basic queries

[Slides 8-16 are essentially covered by Prep4]

Meaning of a query with one relation

```
SELECT name  
FROM Course  
WHERE dept = 'CSC';
```

$$\pi_{\text{name}} (\sigma_{\text{dept}=\text{"csc"}} (\text{Course}))$$

... and with multiple relations

```
SELECT name  
FROM Offering, Took  
WHERE Offering.id = Took.oid and  
      dept = 'CSC';
```

$$\pi_{\text{name}} (\sigma_{\text{Offering.id=Took.id} \wedge \text{dept='csc'}} (\text{Offering} \times \text{Took}))$$

Temporarily renaming a table

- You can rename tables (just for the duration of the statement):

```
SELECT e.name, d.name  
FROM employee e, department d  
WHERE d.name = 'marketing'  
AND e.name = 'Horton';
```

- Can be convenient vs the longer full names:

```
SELECT employee.name, department.name  
FROM employee, department  
WHERE department.name = 'marketing'  
AND employee.name = 'Horton';
```

- This is like ρ in relational algebra.

Self-joins

- As we know, renaming is *required* for self-joins.
- Example:

```
SELECT e1.name, e2.name  
FROM employee e1, employee e2  
WHERE e1.salary < e2.salary;
```

* In SELECT clauses

- A * in the SELECT clause means “all attributes of this relation.”

- Example:

```
SELECT *  
FROM Course  
WHERE dept = 'CSC';
```

Renaming attributes

- Use *AS «new name»* to rename an attribute in the result.

- Example:

```
SELECT name AS title, dept  
FROM Course  
WHERE breadth;
```

Complex Conditions in a WHERE

- We can build boolean expressions with operators that produce boolean results.
 - comparison operators: `=`, `<>`, `<`, `>`, `<=`, `>=`
 - and many other operators:
see section 6.1.2 of the text and chapter 9 of the postgresSQL documentation.
- Note that “not equals” is unusual: `<>`
- We can combine boolean expressions with:
 - Boolean operators: `AND`, `OR`, `NOT`.

Example: Compound condition

- Find 3rd- and 4th-year CSC courses:

```
SELECT *  
FROM Offering  
WHERE dept = 'CSC' AND cnum >= 300;
```

ORDER BY

- To put the tuples in order, add this as the final clause:
`ORDER BY «attribute list» [DESC]`
- The default is ascending order; DESC overrides it to force descending order.
- The attribute list can include expressions: e.g.,
`ORDER BY sales+rentals`
- The ordering is the last thing done before the SELECT, so all attributes are still available.

Case-sensitivity and whitespace

- Example query:

```
SELECT surName  
FROM Student  
WHERE campus = 'StG';
```

- Keywords, like `SELECT`, are not case-sensitive.
 - One convention is to use uppercase for keywords.
- Identifiers, like `Student` are not case-sensitive either.
 - One convention is to use lowercase for attributes, and a leading capital letter followed by lowercase for relations.
- Literal strings, like `'StG'`, are case-sensitive, and require single quotes.
- Whitespace (other than inside quotes) is ignored.

Expressions in SELECT clauses

- Instead of a simple attribute name, you can use an expression in a SELECT clause.
- Operands: attributes, constants
Operators: arithmetic ops, string ops

- Examples:

```
SELECT sid, grade+10 AS adjusted  
FROM Took;
```

```
SELECT dept || cnum  
FROM course;
```

Expressions that are a constant

- Sometimes it makes sense for the whole expression to be a constant (something that doesn't involve any attributes!).

- Example:

```
SELECT dept, cNum,  
       'satisfies' AS breadthRequirement  
FROM Course  
WHERE breadth;
```

Pattern operators

- Two ways to compare a string to a pattern by:
 - «*attribute*» LIKE «*pattern*»
 - «*attribute*» NOT LIKE «*pattern*»
- Pattern is a quoted string
 - % means: any string
 - _ means: any single character
- Example:

```
SELECT *  
FROM Course  
WHERE name LIKE ' %Comp% ' ;
```

Aggregation

Computing on a column

- We often want to compute something across the values in a column.
- `SUM`, `AVG`, `COUNT`, `MIN`, and `MAX` can be applied to a column in a `SELECT` clause.
- Also, `COUNT (*)` counts the number of tuples.
- We call this aggregation.
- Note: To stop duplicates from contributing to the aggregation, use `DISTINCT` inside the brackets.
- **Example:** aggregation.txt

Grouping

- **Example:** group-by.txt
- If we follow a SELECT-FROM-WHERE expression with GROUP BY <attributes>
 - The tuples are grouped according to the values of those attributes, and
 - any aggregation gives us a single value per group.

Restrictions on aggregation

- If any aggregation is used, then each element of the SELECT list must be either:
 - aggregated, or
 - an attribute on the GROUP BY list.
- Otherwise, it doesn't even make sense to include the attribute.

HAVING Clauses

- **Example:** having.txt
- WHERE let's you decide which tuples to keep.
- Similarly, you can decide which *groups* to keep.
- Syntax:
 - . . .
 - GROUP BY «*attributes*»
 - HAVING «*condition*»
- Semantics:

Only groups satisfying the condition are kept.

Requirements on HAVING clauses

- Outside subqueries, HAVING may refer to attributes only if they are either:
 - aggregated, or
 - an attribute on the GROUP BY list.
- (Same requirement as for SELECT clauses with aggregation)

Set operations

Tables can have duplicates in SQL

- A table can have duplicate tuples, unless this would violate an integrity constraint.
- And SELECT-FROM-WHERE statements leave duplicates in unless you say not to.
- Why?
 - Getting rid of duplicates is expensive!
 - We may want the duplicates because they tell us how many times something occurred.

Bags

- SQL treats tables as “bags” (or “multisets”) rather than sets.
- Bags are just like sets, but duplicates are allowed.
- $\{6, 2, 7, 1, 9\}$ is a set (and a bag)
 $\{6, 2, 2, 7, 1, 9\}$ is not a set, but is a bag.
- Like with sets, order doesn't matter.
 $\{6, 2, 7, 1, 9\} = \{1, 2, 6, 7, 9\}$
- **Example:** Tables with duplicates

Union, Intersection, and Difference

- These are expressed as:

(«*subquery*») UNION («*subquery*»)

(«*subquery*») INTERSECT («*subquery*»)

(«*subquery*») EXCEPT («*subquery*»)

- The brackets are mandatory.
- The operands must be queries; you can't simply use a relation name.

Example

```
(SELECT sid
  FROM Took
  WHERE grade > 95)
      UNION
(SELECT sid
  FROM Took
  WHERE grade < 50);
```

Operations \cup , \cap , and $-$ with Bags

- For \cup , \cap , and $-$ the number of occurrences of a tuple in the result requires some thought.
- (But it makes total sense.)

Operations \cup , \cap , and $-$ with Bags

- Suppose tuple t occurs
 - m times in relation R , and
 - n times in relation S .

Operation	Number of occurrences of t in result
$R \cap S$	$\min(m, n)$
$R \cup S$	$m + n$
$R - S$	$\max(m - n, 0)$

Bag vs Set Semantics: which is used

- We saw that a SELECT-FROM-WHERE statement uses bag semantics by default.
 - Duplicates are kept in the result.
- The set operations use set semantics by default.
 - Duplicates are *eliminated* from the result.

Motivation: Efficiency

- When doing projection, it is easier not to eliminate duplicates.
 - Just work one tuple at a time.
- For intersection or difference, it is most efficient to sort the relations first.
 - At that point you may as well eliminate the duplicates anyway.

Controlling Duplicate Elimination

- We can force the result of a SFW query to be a set by using `SELECT DISTINCT ...`
- We can force the result of a set operation to be a bag by using `ALL`, e.g.,

```
(SELECT sid
  FROM Took
 WHERE grade > 95)
      UNION ALL
(SELECT sid
  FROM Took
 WHERE grade < 50);
```

- **Examples:** controlling-dups.txt, except-all.txt

Views

The idea

- A view is a relation defined in terms of stored tables (called base tables) and other views.
- Access a view like any base table.
- Two kinds of view:
 - **Virtual**: no tuples are stored; view is just a query for constructing the relation when needed.
 - **Materialized**: actually constructed and stored.
Expensive to maintain!
- We'll use only virtual views.
 - PostgreSQL did not support materialized views until version 9.3 (which we are not running).

Example: defining a virtual view

- A view for students who earned an 80 or higher in a CSC course.

```
CREATE VIEW toprresults AS
SELECT firstname, surname, cnum
FROM Student, Took, Offering
WHERE
    Student.sid = Took.sid AND
    Took.oid = Offering.oid AND
    grade >= 80 AND dept = 'CSC';
```

Uses for views

- Break down a large query.
- Provide another way of looking at the same data, e.g., for one category of user.