

# Outer Joins

# The joins you know from RA

These can go in a FROM clause:

Expression	Meaning
$R, S$	$R \times S$
$R \text{ cross join } S$	
$R \text{ natural join } S$	$R \bowtie S$
$R \text{ join } S \text{ on Condition}$	$R \bowtie_{\text{condition}} S$

# In practise, natural join is brittle

- A working query can be broken by adding a column to a schema.
  - Example:

```
SELECT sid, instructor
FROM Student NATURAL JOIN Took
      NATURAL JOIN Offering;
```
  - What if we add a column called `campus` to `Offering`?
- Also, having implicit comparisons impairs readability.
- Best practise: Don't use natural join.

# Dangling tuples

- With joins that require some attributes to match, tuples lacking a match are left out of the results.
- We say that they are “dangling”.
- An **outer join** preserves dangling tuples by padding them with **NULL** in the other relation.
- A join that doesn't pad with **NULL** is called an **inner join**.

# Three kinds of outer join

- **LEFT** OUTER JOIN
  - Preserves dangling tuples from the relation on the LHS by padding with nulls on the RHS.
- **RIGHT** OUTER JOIN
  - The reverse.
- **FULL** OUTER JOIN
  - Does both.

# Example: joining R and S various ways

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL JOIN S

A	B	C
1	2	3

# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL FULL JOIN S

A	B	C
1	2	3
4	5	NULL
NULL	6	7

# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL LEFT JOIN S

A	B	C
1	2	3
4	5	NULL



# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL RIGHT JOIN S

A	B	C
1	2	3
NULL	6	7

# Summary of join expressions

## Cartesian product

`A CROSS JOIN B`

same as `A, B`

## Theta-join

`A JOIN B ON C`

✓ `A {LEFT|RIGHT|FULL} JOIN B ON C`

## Natural join

`A NATURAL JOIN B`

✓ `A NATURAL {LEFT|RIGHT|FULL} JOIN B`

✓ indicates that tuples are padded when needed.

# Keywords INNER and OUTER

- There are keywords `INNER` and `OUTER`, but you never need to use them.
- Your intentions are clear anyway:
  - You get an outer join iff you use the keywords `LEFT`, `RIGHT`, or `FULL`.
  - If you don't use the keywords `LEFT`, `RIGHT`, or `FULL` you get an inner join.

# Impact of having null values

# Missing Information

- Two common scenarios:
  - Missing value.  
E.g., we know a student has some email address, but we don't know what it is.
  - Inapplicable attribute.  
E.g., the value of attribute spouse for an unmarried person.

# Representing missing information

- One possibility: use a special value as a placeholder. E.g.,
  - If age unknown, use 0.
  - If StNum unknown, use 999999999.
- Implications?
- Better solution: use a value not in any domain. We call this a null value.
- Tuples in SQL relations can have **NULL** as a value for one or more components.

# Checking for null values

- You can compare an attribute value to **NULL** with
  - **IS NULL**
  - **IS NOT NULL**
- **Example:**

```
SELECT *  
FROM Course  
WHERE breadth IS NULL;
```

# In SQL we have 3 truth-values

- Because of **NULL**, we need three truth-values:
  - If one or both operands to a comparison is **NULL**, the comparison *always* evaluates to **UNKNOWN**.
  - Otherwise, comparisons evaluate to **TRUE** or **FALSE**.



# Combining truth values

- We need to know how the three truth-values combine with **AND**, **OR** and **NOT**.
- Can think of it in terms of the truth table.
- Or can think in terms of numbers:
  - **TRUE** = 1, **FALSE** = 0, **UNKNOWN** = 0.5
  - **AND** is min, **OR** is max,
  - **NOT** x is  $(1-x)$ , i.e., it “flips” the value

# The three-valued truth table

A	B	A and B	A or B
T	T	T	T
TF or FT		F	T
F	F	F	F
TU or UT		U	T
FU or UF		F	U
U	U	U	U

# Thinking of the truth-values as numbers

A	B	as nums	A and B	min	A or B	max
T	T	1, 1	T	1	T	1
TF or FT		1, 0	F	0	T	1
F	F	0, 0	F	0	F	0
TU or UT		1, 0.5	U	0.5	T	1
FU or UF		0, 0.5	F	0	U	0.5
U	U	0.5, 0.5	U	0.5	U	0.5

# Thinking of the truth-values as

A	as a num, x	not A	1 - x
T	1	F	0
F	0	T	1
U	0.5	U	0.5

# Surprises from 3-valued logic

- Some laws you are used to still hold in three-valued logic. For example,
  - **AND** is commutative.
- But others don't. For example,
  - The law of the excluded middle breaks:  
 $(p \text{ or } (\text{NOT } p))$  might not be **TRUE**!
  - $(0 * x)$  might not be 0.

# Impact of null values on WHERE

- A tuple is in a query result iff the WHERE clause is **TRUE**.
- **UNKNOWN** is not good enough.
- “WHERE is picky.”
- Example: **where-null**

# Impact of null values on aggregation

- Summary: Aggregation ignores **NULL**.
  - **NULL** never contributes to a sum, average, or count, and
  - can never be the minimum or maximum of a column (unless every value is **NULL**).
- If there are no *non-NULL* values in a column, then the result of the aggregation is **NULL**.
  - Exception: **COUNT** of an empty set is 0.

# Aggregation ignores nulls

	some nulls in A	All nulls in A
<code>min(A)</code>	ignore the nulls	null
<code>max(A)</code>		
<code>sum(A)</code>		
<code>avg(A)</code>		
<code>count(A)</code>		0
<code>count(*)</code>	all tuples count	

 Example: aggregation-nulls



# More re the impact of null values

- Other corner cases to think about:
  - `SELECT DISTINCT`: are 2 `NULL` values equal?
  - natural join: are 2 `NULL` values equal?
  - set operations: are 2 `NULL` values equal?
- And later, when we learn about constraints:
  - `UNIQUE` constraint: do 2 `NULL` values violate?
- This behaviour may vary across DBMSs.

# Summary re NULL

- Any comparison with NULL yields UNKNOWN.
- WHERE is picky: it only accepts TRUE.
- Therefore NATURAL JOIN is picky too.
- Aggregation ignores NULL.
- In other situations where NULLs matter
  - when a truth-value may be NULL
  - when it matters whether two NULL are considered the same

Don't assume. Behaviour may vary by DBMS.