

# SQL Exercises: Subqueries

## Schema

Student( <u>sID</u> , surName, firstName, campus, email, cgpa)	Offering[dept, cNum] $\subseteq$ Course[dept, cNum]
Course( <u>dept</u> , <u>cNum</u> , name, breadth)	Took[sID] $\subseteq$ Student[sID]
Offering( <u>oID</u> , dept, cNum, term, instructor)	Took[oID] $\subseteq$ Offering[oID]
Took( <u>sID</u> , <u>oID</u> , grade)	

## Questions

1. What does this query do? (The || operator concatenates two strings.)

```
SELECT sid, dept || cnum as course, grade
FROM Took,
    (SELECT *
     FROM Offering
     WHERE instructor = 'Horton') Hofferings
WHERE Took.oid = Hofferings.oid;
```

2. What does this query do?

```
SELECT sid, surname
FROM Student
WHERE cgpa >
    (SELECT cgpa
     FROM Student
     WHERE sid = 99999);
```

3. What does this query do?

```
SELECT sid, dept || cnum AS course, grade
FROM Took JOIN Offering ON Took.oid = Offering.oid
WHERE
    grade >= 80 AND
    (cnum, dept) IN (
        SELECT cnum, dept
        FROM Took JOIN Offering ON Took.oid = Offering.oid
        JOIN Student ON Took.sid = Student.sid
        WHERE surname = 'Lakemeyer');
```

4. (a) Suppose we have these relations: R(a, b) and S(b, c). What does this query do?

```
SELECT a
FROM R
WHERE b in (SELECT b FROM S);
```

- (b) Can we express this query without using subqueries?

5. What does this query do?

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
    SELECT *
    FROM Offering
    WHERE
        oid <> Off1.oid AND
        instructor = Off1.instructor );
```

6. What does this query do?

```
SELECT DISTINCT oid
FROM Took
WHERE EXISTS (
    SELECT *
    FROM Took t, Offering o
    WHERE
        t.oid = o.oid AND
        t.oid <> Took.oid AND
        o.dept = 'CSC' AND
        took.sid = t.sid )
ORDER BY oid;
```

7. Now let's write some queries! For each course find the instructor who has taught the most offerings of it. If there are ties, include them all Report the course (eg "csc343"), instructor and the number of offerings of the course by that instructor. Suggestion: Use one or more views to hold intermediate step(s).

**Solution:**

```
-- This intermediate result is helpful:
CREATE VIEW Counts as
SELECT dept || cnum as course, instructor, count(oid)
FROM Offering
GROUP BY cnum, dept, instructor;

-- Let's take a look at what this computes.
-- (Our dataset doesn't give this view a very good test.)
SELECT * from Counts;
```

course	instructor	count
CSC148	Miller	1
CSC148	Jepson	2
EEB263	Suzuki	1
CSC343	Mylopoulos	2
EEB216	Suzuki	1
ENG235	Richler	1
ENV200	Suzuki	1
EEB263	Johancsik	1
ENG235	Percy	1
HIS220	Dow	1
CSC343	Horton	1

CSC148		Chechik		1
EEB150		Mendel		1
CSC343		Truta		1
ENV320		Suzuki		1
ENG205		Reisman		1
HIS220		Young		1
ENG205		Atwood		1
CSC263		Horton		2
ENG110		Atwood		1
HIS296		Young		1
CSC207		Gries		2
ANT200		Zorich		1
ANT203		Davies		1
ENG110		Percy		1
ANT203		Zorich		1
CSC343		Heap		1
CSC320		Jepson		2
CSC207		Craig		2
CSC263		Craig		1

(30 rows)

-- Now we can solve the problem using a subquery:

```
SELECT course, instructor, count
FROM Counts C1
WHERE count >= ALL (
    SELECT count
    FROM Counts C2
    WHERE C1.course = C2.course )
ORDER BY C1.course;
```

-- Here's another version:

```
SELECT course, instructor, count
FROM Counts C1
WHERE count = (
    SELECT max(count)
    FROM Counts C2
    WHERE C1.course = C2.course )
ORDER BY C1.course;
```

-- Here's what they both produce:

course		instructor		count
-----	+	-----	+	-----
ANT200		Zorich		1
ANT203		Zorich		1
ANT203		Davies		1
CSC148		Jepson		2
CSC207		Craig		2
CSC207		Gries		2
CSC263		Horton		2
CSC320		Jepson		2
CSC343		Mylopoulos		2
EEB150		Mendel		1

EEB216	Suzuki	1
EEB263	Suzuki	1
EEB263	Johancsik	1
ENG110	Atwood	1
ENG110	Percy	1
ENG205	Atwood	1
ENG205	Reisman	1
ENG235	Richler	1
ENG235	Percy	1
ENV200	Suzuki	1
ENV320	Suzuki	1
HIS220	Dow	1
HIS220	Young	1
HIS296	Young	1

(24 rows)

8. Let's say that a course has level "junior" if its cNum is between 100 and 299 inclusive, and has level "senior" if its cNum is between 300 and 499 inclusive. Report the average grade, across all departments and course offerings, for all junior courses and for all senior courses. Report your answer in a table that looks like this:

level		levelavg
-----		-----
junior		
senior		

Each average should be an average of the individual student grades, not an average of the course averages.

**Solution:**

```
CREATE VIEW Grades AS
SELECT cnum, dept, grade
FROM Offering natural join Took;

(SELECT 'junior' AS level, avg(grade) AS levelavg
FROM Grades
WHERE cnum >= 100 AND cnum <= 299)
union
(SELECT 'senior' AS level, avg(grade) AS levelavg
FROM Grades
WHERE cnum >= 300 AND cnum <= 499);
```

level		levelavg
-----	+	-----
junior		75.0952380952380952
senior		77.5000000000000000

(2 rows)