

-- Recap of GROUP BY

dbsrv1% psql csc343h-dianeh  
psql (9.1.14)  
Type "help" for help.

csc343h-dianeh=> set search\_path to university;  
SET  
csc343h-dianeh=> \?  
csc343h-dianeh=> \d

          List of relations  
  Schema | Name | Type | Owner  
-----+-----+-----+-----  
university | course | table | dianeh  
university | offering | table | dianeh  
university | student | table | dianeh  
university | took | table | dianeh  
(4 rows)

-- We saw queries like these last class.

-- Try them yourself to get more comfortable with what they do.

csc343h-dianeh=> select \* from took;  
csc343h-dianeh=> select \* from took order by oid;  
csc343h-dianeh=> select \* from took group by oid;  
ERROR: column "took.sid" must appear in the GROUP BY clause or be used in an aggregate function  
LINE 1: select \* from took group by oid;  
          ^

csc343h-dianeh=> select oid, max(sid), avg(grade) from took group by oid;  
csc343h-dianeh=> select oid, max(sid), avg(grade) from took group by oid order by avg(grade);  
csc343h-dianeh=> select oid, max(sid), avg(grade) from took group by oid order by avg(grade) desc;  
csc343h-dianeh=> select oid, max(sid), avg(grade) from took group by oid order by oid;  
csc343h-dianeh=> select oid, max(sid), avg(grade) from took group by oid order by avg(grade), oid;

-- Worksheet

csc343h-dianeh=> -- Question 1  
csc343h-dianeh=> -- First let's find these values overall.  
csc343h-dianeh=> select avg(grade), min(grade), max(grade) from took;  
          avg | min | max  
-----+-----+-----  
75.6296296296296296 | 0 | 100  
(1 row)

csc343h-dianeh=> -- Now let's find them per oid.

csc343h-dianeh=> select avg(grade), min(grade), max(grade) from took group by oid;  
          avg | min | max  
-----+-----+-----  
59.0000000000000000 | 39 | 98  
60.6666666666666667 | 45 | 75  
70.5000000000000000 | 52 | 89  
92.0000000000000000 | 91 | 93  
69.5000000000000000 | 46 | 94  
91.0000000000000000 | 91 | 91  
87.2500000000000000 | 79 | 99  
31.0000000000000000 | 0 | 62  
71.0000000000000000 | 71 | 71  
79.0000000000000000 | 39 | 99  
92.0000000000000000 | 92 | 92  
73.5000000000000000 | 17 | 100  
97.0000000000000000 | 97 | 97  
74.0000000000000000 | 72 | 78  
71.0000000000000000 | 71 | 71  
82.0000000000000000 | 82 | 82

78.0000000000000000	70	82
75.0000000000000000	75	75
74.6666666666666667	59	89
95.6666666666666667	90	99
75.0000000000000000	54	96
78.0000000000000000	78	78
83.0000000000000000	71	89

(23 rows)

csc343h-dianeH=> -- The results might be more interesting if we added in the oid, and ordered them somehow.

csc343h-dianeH=> -- Let's do it by average grade.

csc343h-dianeH=> select oid, avg(grade), min(grade), max(grade) from took group by oid order by avg(grade);

oid	avg	min	max
15	31.0000000000000000	0	62
14	59.0000000000000000	39	98
34	60.6666666666666667	45	75
17	69.5000000000000000	46	94
27	70.5000000000000000	52	89
21	71.0000000000000000	71	71
26	71.0000000000000000	71	71
16	73.5000000000000000	17	100
6	74.0000000000000000	72	78
5	74.6666666666666667	59	89
35	75.0000000000000000	75	75
22	75.0000000000000000	54	96
31	78.0000000000000000	70	82
9	78.0000000000000000	78	78
11	79.0000000000000000	39	99
3	82.0000000000000000	82	82
7	83.0000000000000000	71	89
1	87.2500000000000000	79	99
28	91.0000000000000000	91	91
8	92.0000000000000000	91	93
38	92.0000000000000000	92	92
13	95.6666666666666667	90	99
39	97.0000000000000000	97	97

(23 rows)

csc343h-dianeH=> -- Question 2:

csc343h-dianeH=> -- Which can go in the select?:

csc343h-dianeH=> -- (Notice that this query joins 2 tables.)

csc343h-dianeH=> -- sid: no

csc343h-dianeH=> -- count(sid): yes

csc343h-dianeH=> -- grade: no

csc343h-dianeH=> -- avg(grade): yes

csc343h-dianeH=> -- dept: yes -- there is one dept per dept, so to speak

csc343h-dianeH=> -- count(dept): yes

csc343h-dianeH=> -- term: no

csc343h-dianeH=> -- min(term): yes

csc343h-dianeH=>

csc343h-dianeH=> -- Let's run the query from Question 2:

csc343h-dianeH=> select count(sid), avg(grade), dept, count(dept), min(term)

csc343h-dianeH-> from offering, took

csc343h-dianeH-> where offering.oid = took.oid

csc343h-dianeH-> group by dept;

count	avg	dept	count	min
4	69.5000000000000000	ENV	4	20089
6	78.1666666666666667	EEB	6	20081
8	78.5000000000000000	ANT	8	20081
1	97.0000000000000000	HIS	1	20081
24	79.6666666666666667	CSC	24	20081

11 | 63.63636363636364 | ENG | 11 | 20081  
(6 rows)

```
csc343h-dianehe=> -- What if we include sid?
csc343h-dianehe=> -- It would make no sense, since there is not one sid per dept.
csc343h-dianehe=> -- So we get that error message that, by now, makes sense too.
csc343h-dianehe=> select count(sid), avg(grade), dept, count(dept), min(term), sid
csc343h-dianehe-> from offering, took
csc343h-dianehe-> where offering.oid = took.oid
csc343h-dianehe-> group by dept;
ERROR: column "took.sid" must appear in the GROUP BY clause or be used in an aggregate function
LINE 1: ...ct count(sid), avg(grade), dept, count(dept), min(term), sid
```

-- Next feature: HAVING

```
csc343h-dianehe=> -- Here's a plain GROUP BY.
csc343h-dianehe=> select oid, avg(grade), count(*)
csc343h-dianehe-> from took
csc343h-dianehe-> group by oid;
```

oid	avg	count
14	59.000000000000000	3
34	60.666666666666667	3
27	70.500000000000000	2
8	92.000000000000000	2
17	69.500000000000000	4
28	91.000000000000000	1
1	87.250000000000000	4
15	31.000000000000000	2
26	71.000000000000000	1
11	79.000000000000000	4
38	92.000000000000000	1
16	73.500000000000000	4
39	97.000000000000000	1
6	74.000000000000000	3
21	71.000000000000000	1
3	82.000000000000000	1
31	78.000000000000000	4
35	75.000000000000000	1
5	74.666666666666667	3
13	95.666666666666667	3
22	75.000000000000000	2
9	78.000000000000000	1
7	83.000000000000000	3

(23 rows)

```
csc343h-dianehe=> -- Now let's filter the results to only the groups having a desired property.
csc343h-dianehe=> select oid, avg(grade), count(*)
csc343h-dianehe-> from took
csc343h-dianehe-> group by oid
csc343h-dianehe-> having count(*) > 1;
```

oid	avg	count
14	59.000000000000000	3
34	60.666666666666667	3
27	70.500000000000000	2
8	92.000000000000000	2
17	69.500000000000000	4
1	87.250000000000000	4
15	31.000000000000000	2
11	79.000000000000000	4
16	73.500000000000000	4
6	74.000000000000000	3
31	78.000000000000000	4

5	74.666666666666667	3
13	95.666666666666667	3
22	75.000000000000000	2
7	83.000000000000000	3

(15 rows)

```
csc343h-dianeh=> -- Here's the same query, but with a filter that involves an unaggregated attribute.
csc343h-dianeh=> select oid, avg(grade), count(*)
csc343h-dianeh-> from took
csc343h-dianeh-> group by oid
csc343h-dianeh-> having oid > 10;
```

oid	avg	count
14	59.000000000000000	3
34	60.666666666666667	3
27	70.500000000000000	2
17	69.500000000000000	4
28	91.000000000000000	1
15	31.000000000000000	2
26	71.000000000000000	1
11	79.000000000000000	4
38	92.000000000000000	1
16	73.500000000000000	4
39	97.000000000000000	1
21	71.000000000000000	1
31	78.000000000000000	4
35	75.000000000000000	1
13	95.666666666666667	3
22	75.000000000000000	2

(16 rows)

```
csc343h-dianeh=> -- But if we filter using grade unaggregated, it doesn't work.
csc343h-dianeh=> -- That makes complete sense.
csc343h-dianeh=> -- We have grouped by oid, so there will be one row per oid (if that oid passes
csc343h-dianeh=> -- the filter). But the filter is grade < 50, and there is not one grade to check
csc343h-dianeh=> -- per oid.
csc343h-dianeh=> select oid, avg(grade), count(*)
csc343h-dianeh-> from took
csc343h-dianeh-> group by oid
csc343h-dianeh-> having grade < 50;
ERROR: column "took.grade" must appear in the GROUP BY clause or be used in an aggregate function
LINE 4: having grade < 50;
              ^
```

```
csc343h-dianeh=> -- If we change the filter to min(grade) < 50, the query makes sense again. There
csc343h-dianeh=> -- is only one average grade per oid, so we can compare it to 50.
csc343h-dianeh=> select oid, min(grade), avg(grade), count(*)
csc343h-dianeh-> from took
csc343h-dianeh-> group by oid
csc343h-dianeh-> having min(grade) < 50;
```

oid	min	avg	count
14	39	59.000000000000000	3
34	45	60.666666666666667	3
17	46	69.500000000000000	4
15	0	31.000000000000000	2
11	39	79.000000000000000	4
16	17	73.500000000000000	4

(6 rows)

```
csc343h-dianeh=> -- We can even filter on something that is not in the SELECT clause.
csc343h-dianeh=> select oid, avg(grade), count(*)
csc343h-dianeh-> from took
csc343h-dianeh-> group by oid
csc343h-dianeh-> having min(grade) < 50;
```

oid	avg	count
14	59.000000000000000	3
34	60.666666666666667	3
17	69.500000000000000	4
15	31.000000000000000	2
11	79.000000000000000	4
16	73.500000000000000	4

(6 rows)

-- Back to the worksheet

csc343h-dianehe=> -- Question 3 (as originally published on the sheet)

```
csc343h-dianehe=>
csc343h-dianehe=> select sid, avg(grade)
csc343h-dianehe=> from took
csc343h-dianehe=> group by sid
csc343h-dianehe=> having sid > 22222;
```

sid	avg
99132	76.2857142857142857
99999	84.583333333333333
98000	83.200000000000000

(3 rows)

csc343h-dianehe=> -- Here it is without the filter, so we can see what was removed.

```
csc343h-dianehe=> select sid, avg(grade)
csc343h-dianehe=> from took
csc343h-dianehe=> group by sid;
```

sid	avg
11111	29.600000000000000
98000	83.200000000000000
99132	76.2857142857142857
99999	84.583333333333333
157	75.933333333333333

(5 rows)

csc343h-dianehe=> -- Question 3 (the new version we did in the noon and 2pm sections).

```
csc343h-dianehe=> -- For each student who has passed at least 2 courses,
csc343h-dianehe=> -- report their sid and average grade on the courses that they passed.
```

```
csc343h-dianehe=> select sid, avg(grade)
csc343h-dianehe=> from took
csc343h-dianehe=> where grade >= 50
csc343h-dianehe=> group by sid
csc343h-dianehe=> having count(*) >= 2;
```

sid	avg
98000	83.200000000000000
99132	82.500000000000000
99999	84.583333333333333
157	78.5714285714285714

(4 rows)

csc343h-dianehe=> -- Let's put the count in so that we can see what was filtered.

```
csc343h-dianehe=> -- Oh, it filtered nothing because everyone has passed at least two
csc343h-dianehe=> -- courses.
```

```
csc343h-dianehe=> select sid, avg(grade), count(*)
csc343h-dianehe=> from took
csc343h-dianehe=> where grade >= 50
csc343h-dianehe=> group by sid;
```

sid	avg	count
-----	-----	-------

```

98000 | 83.2000000000000000 | 15
99132 | 82.5000000000000000 | 6
99999 | 84.5833333333333333 | 12
157 | 78.5714285714285714 | 14
(4 rows)

```

```

csc343h-dianeH=> -- Let's make the condition more stringent,
csc343h-dianeH=> -- so that there will be some actual filtering.

```

```

csc343h-dianeH=> select sid, avg(grade), count(*)
csc343h-dianeH-> from took
csc343h-dianeH-> where grade >= 50
csc343h-dianeH-> group by sid
csc343h-dianeH-> having count(*) >= 10;
  sid |          avg          | count
-----+-----+-----
 98000 | 83.2000000000000000 |    15
 99999 | 84.5833333333333333 |    12
   157 | 78.5714285714285714 |    14
(3 rows)

```

```

csc343h-dianeH=> -- Question 4:
csc343h-dianeH=> select sid
csc343h-dianeH-> from took
csc343h-dianeH-> group by sid
csc343h-dianeH-> having avg(grade) > 80;
  sid
-----
 98000
 99999
(2 rows)

```

```

csc343h-dianeH=> -- Let's put the average grade in so we can check our results.
csc343h-dianeH=> select sid, avg(grade)
csc343h-dianeH-> from took
csc343h-dianeH-> group by sid
csc343h-dianeH-> having avg(grade) > 80;
  sid |          avg          |
-----+-----+-----
 98000 | 83.2000000000000000 |
 99999 | 84.5833333333333333 |
(2 rows)

```