

== Learning XPath through Examples ==

= Basics =

General idea of paths is familiar from unix

- `ls course-*/assignments/a1/solution/*.py`

Simplest example

- `fn:doc("quiz.xml")/quiz/questions/mc-question/question`
 - just a straight-up path
 - notice that you get a list of results
 - comma-separated
 - they are in the order in which they appear in the doc
 - it is case sensitive: break it and see

Going up the tree

- modify it to go one node less deep and see result
- Aside: what if you create a path that doesn't exist
 - eg, take out a step
 - get the empty list
- keep shortening the path until the path is empty:
 - you get the document node

`fn:doc("quiz.xml")/quiz` vs `fn:doc("quiz.xml")`

- `fn:doc("quiz.xml")/quiz`
 - a sequence of one item
 - that item is the node for the element that is the document's root (quiz), and contains the whole document within it
- `fn:doc("quiz.xml")`
 - a sequence of one item
 - that item is the document node for the whole file
- the difference is clear if you go back to the picture of the tree that the xml represents [see the "XML & DTD" lecture slides]

Using the `text()` function

- if you add `text()` at the end, you get the body of the element, not the element.
 - eg where the element contains #PCDATA
 - `fn:doc("quiz.xml")/quiz/questions/mc-question/question/text()`
 - get text of the body of the question elements
 - eg where it doesn't
 - take out "questions"
 - the text part of an mc-question element is just some whitespace
 - notice inconsistencies in the whitespace
 - quiz-a.xml has the whitespace manipulated in interesting ways.
 - try running the same query on quiz-a.xml to see the effects.
- `fn:doc("quiz-a.xml")/quiz/questions/mc-question/question/text()`

Go to an attribute

- `@att-name` (and you do need a slash before it)
- `fn:doc("quiz.xml")/quiz/questions/mc-question/@solution`
- get primitive-type result, not node-type
 - attribute solution {"1"}, attribute solution {"4"}, attribute solution {"3"}
- another example:
`fn:doc("quiz.xml")/quiz/class-responses/student/@sid`
 - attribute sid {"s998801234"},
 - attribute sid {"s001078452"},
 - attribute sid {"s997733991"},
 - attribute sid {"s555555555"}

- can't go further on the path because all attribute nodes are leaves in the document tree

= More features =

Start at any blah-node

- begin path with //blah
- meaning: begin by finding any blah node anywhere at root or lower
- Saves a lot of typing!
Also handy if you can't remember or don't want to look up a long path to a node.
- fn:doc("quiz.xml")//question
Notice that a question node can occur in two different contexts.
- fn:doc("quiz.xml")//@sid
Here it's an attribute we're finding
- fn:doc("quiz.xml")//@qid
Why do we see the same qid values repeated so many times?
- can actually put // anywhere in the path expression
fleep...floop//blah
Start out following the path expression to floop, then from there find a blah-node ANYWHERE underneath
- Can continue the path expression afterward
- fn:doc("quiz.xml")//mc-question/@qid
qid values but only in an mc-question
- fn:doc("quiz.xml")//questions//@qid
qid values but only under some question node
- fn:doc("quiz.xml")/quiz/class-responses//@qid
qid values but only somewhere under a class-responses node

Wildcard

- fn:doc("quiz.xml")/*/@qid
- This yields nothing!
- The reason is that * represents any one tag.
- fn:doc("quiz.xml")/*/*/*/@qid
- This will only match a qid attribute that is at level 4.
- fn:doc("quiz.xml")/*/*/*/*/@qid
- This will only match a qid attribute that is at level 5.

Getting heterogeneous results

- Eg:
fn:doc("quiz.xml")//questions/*/*
- If you have an element e with different kinds of subelement, path expression
.../e/*
will give you results that include those two different kinds of subelement.
- This is one way to get different kinds of items in the resulting output.
- here we get question elements and option elements
And notice that they come from different kinds of contexts (mc-question and tf-question).

Indexing into a node set

- Eg:
fn:doc("quiz.xml")//class-responses/student
fn:doc("quiz.xml")//class-responses/student[1]
- note that it counts from 1, not 0.
- and if the index isn't valid?
fn:doc("quiz.xml")//class-responses/student[22]

= Selection conditions =

Selection conditions

- prunes the paths that have been found
- Append this to the end of the path expression:
[expression]

Examples with attributes

- within the []-ed expression, use
@att-name to refer to the value of an attribute of the current element
- Go down to mc-question and only keep those whose qid attribute is >= "Q555"
fn:doc("quiz.xml")//mc-question[@qid >= "Q555"]
- One with a compound condition
fn:doc("quiz.xml")//response[@qid="Q888" and @answer != "4"]

Examples with elements

- within the []-ed expression, use
subelementName to refer to a subelement of the current element
"." to refer to the entirety of the current element
- Go down to a race element, and
only keep those with a result element with value <= 3.50
fn:doc("races.xml")/races/race[result <= 3.50]
- It's existentially quantified: One or more such race elements will do
- What you get back is race elements that satisfy the condition
- Go down to result elements and
only keep those that themselves (.) have value <= 3.50
fn:doc("races.xml")/races/race/result[. <= 3.50]
- What you get back is result elements that satisfy the condition

Can compare two attributes/elements -- not just compare to constants

- Here we compare an attribute and an element
fn:doc("races.xml")/races/race[@name = sponsor]
- in this dtd, can even compare the same element to itself
fn:doc("races.xml")/races/race[result > result]
- It's existentially quantified: One or pairs of such result elements will do
- Notice it doesn't report DawnAtTheDon's,
because there are no two results where one is > the other.

Can continue the path expression beyond the selection

- Example: Find races with names in a certain range,
then pull out any result below
fn:doc("races.xml")/races/*[@name > "K"]/result
- Notice that we get every result element that's under one of the elements
found by
fn:doc("races.xml")/races/*[@name > "K"]

= Cool functions =

- See <http://www.w3.org/TR/xpath-functions>
- average etc:
fn:doc("races.xml")/races/race/avg(result)
fn:doc("races.xml")/races/race/max(result)
fn:doc("races.xml")/races/race/count(result)
Why do we get multiple answers??
For each *race* element, we are asking for a count of
its result sub-elements.
(There is only one race element.)
- A slightly different use of count:
fn:doc("races.xml")/races/count(race)
Why do we now get one answer??
For each *races* element, we are asking for a count of

its race sub-selements.

(There are several races elements.)

- Using these functions in a condition
fn:doc("races.xml")/races/race[avg(result)>4.50]
fn:doc("races.xml")/races/race[avg(result)>4.50 or min(result)>2.50]
- Then can continue the path expression
fn:doc("races.xml")/races/race[avg(result)>4.50 or min(result)>2.50]/result
- To get the overall average: